

Adaptive Communication-Aware Pipeline Parallelism for Multimodal Transformers: A Proof-of-Concept Study

Distributed Deep Learning Course Project

Taught by Zhao Zhang

Taqiya Ehsan, Khizar Anjum
{te137, ma1629}@rutgers.edu

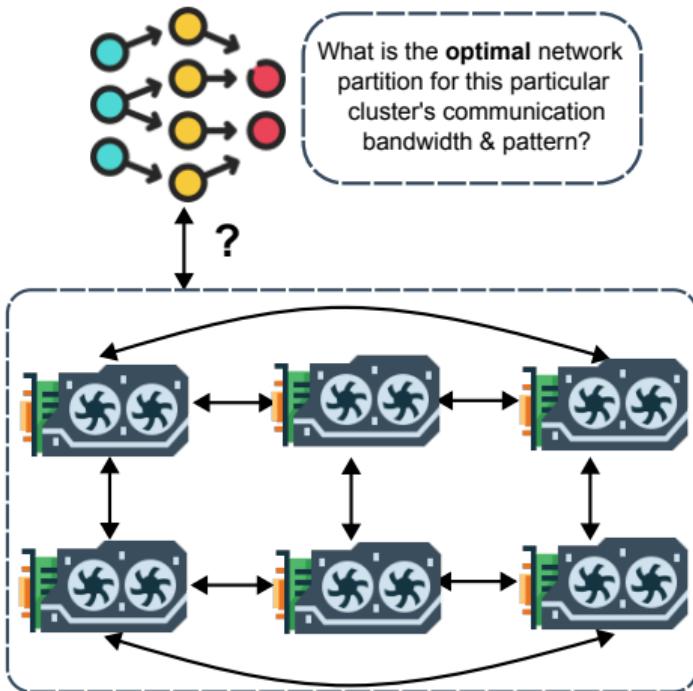
Rutgers University
Department of Electrical and Computer Engineering

Fall 2025

Original Vision: Adaptive Pipeline Parallelism for Multimodal Transformers

Initial Goal:

- **Adaptive pipeline parallelism system** for multimodal vision-language models (BLIP)



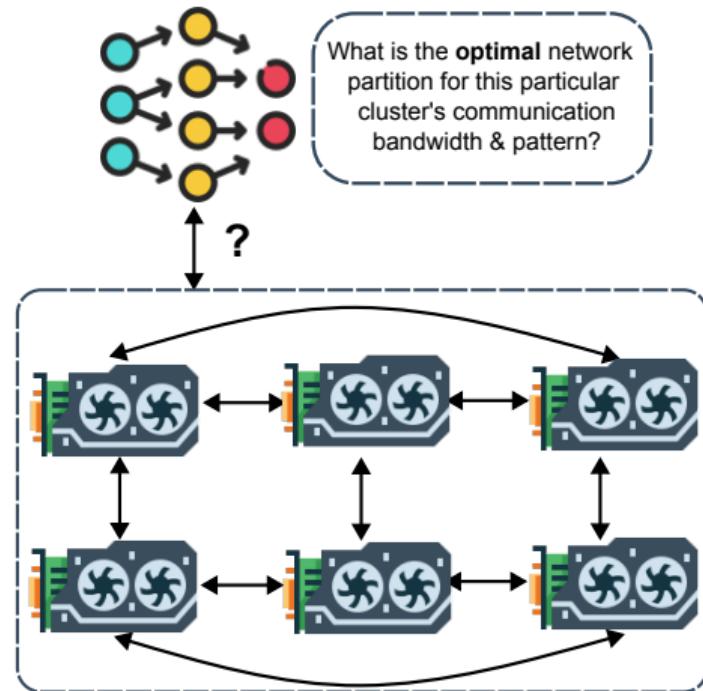
^aM. Shoeybi et al., "Megatron-Lm: Training multi-billion parameter language models using model parallelism," *arXiv preprint arXiv:1909.08053*, 2019 .

^bD. Narayanan et al., "Efficient large-scale language model training on gpu clusters using megatron-Lm," in *ACM/IEEE SC*, 2021.

Original Vision: Adaptive Pipeline Parallelism for Multimodal Transformers

Initial Goal:

- **Adaptive pipeline parallelism system** for multimodal vision-language models (BLIP)
- **Dynamic adjustment** of pipeline configurations based on:
 - Real-time network conditions
 - Heterogeneous compute patterns



^aM. Shoeybi et al., "Megatron-Lm: Training multi-billion parameter language models using model parallelism," *arXiv preprint arXiv:1909.08053*, 2019 .

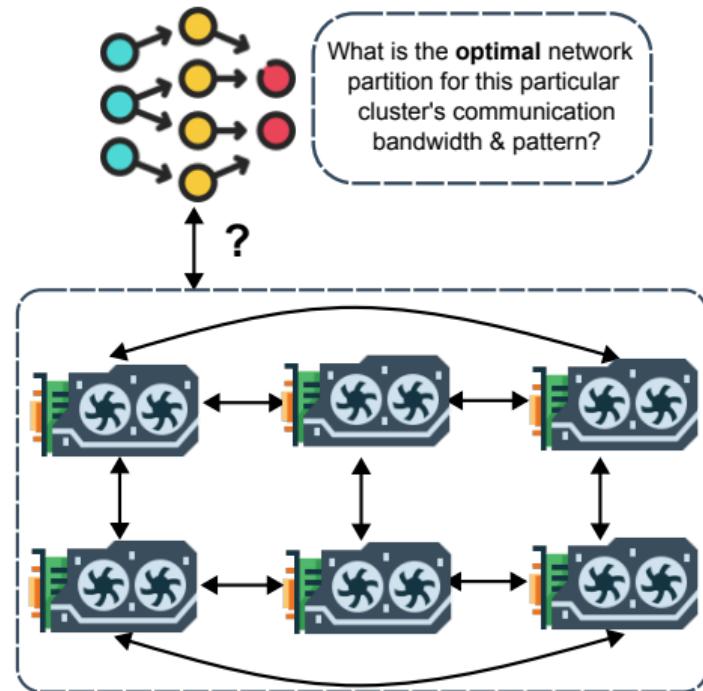
^bD. Narayanan et al., "Efficient large-scale language model training on gpu clusters using megatron-Lm," in *ACM/IEEE SC*, 2021.

Figure: Envisioned adaptive system responding to network conditions

Original Vision: Adaptive Pipeline Parallelism for Multimodal Transformers

Initial Goal:

- **Adaptive pipeline parallelism system** for multimodal vision-language models (BLIP)
- **Dynamic adjustment** of pipeline configurations based on:
 - Real-time network conditions
 - Heterogeneous compute patterns
- Address inefficiencies in **static** pipeline implementations



^aM. Shoeybi et al., "Megatron-Lm: Training multi-billion parameter language models using model parallelism," *arXiv preprint arXiv:1909.08053*, 2019 .

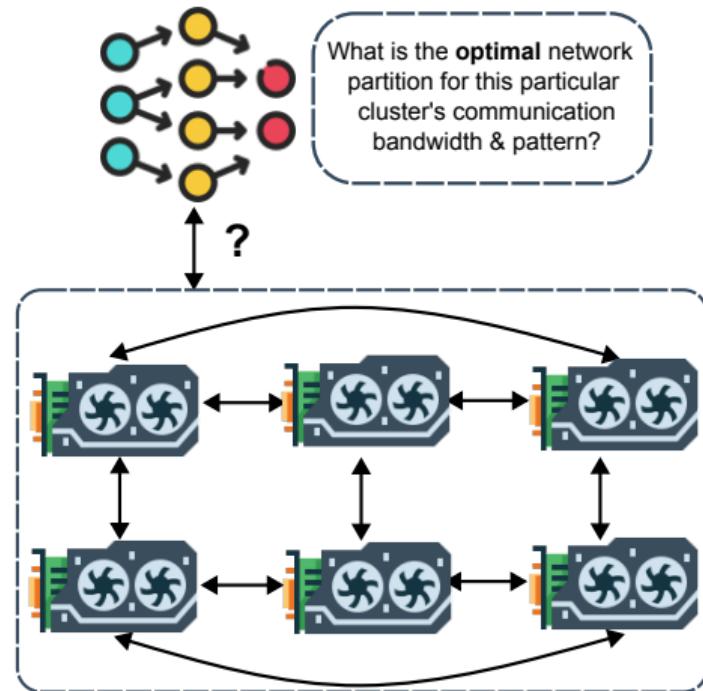
^bD. Narayanan et al., "Efficient large-scale language model training on gpu clusters using megatron-Lm," in *ACM/IEEE SC*, 2021.

Figure: Envisioned adaptive system responding to network conditions

Original Vision: Adaptive Pipeline Parallelism for Multimodal Transformers

Initial Goal:

- **Adaptive pipeline parallelism system** for multimodal vision-language models (BLIP)
- **Dynamic adjustment** of pipeline configurations based on:
 - Real-time network conditions
 - Heterogeneous compute patterns
- Address inefficiencies in **static** pipeline implementations
- Extends NVIDIA's Megatron-LM framework^{a,b}



^aM. Shoeybi et al., "Megatron-Lm: Training multi-billion parameter language models using model parallelism," *arXiv preprint arXiv:1909.08053*, 2019 .

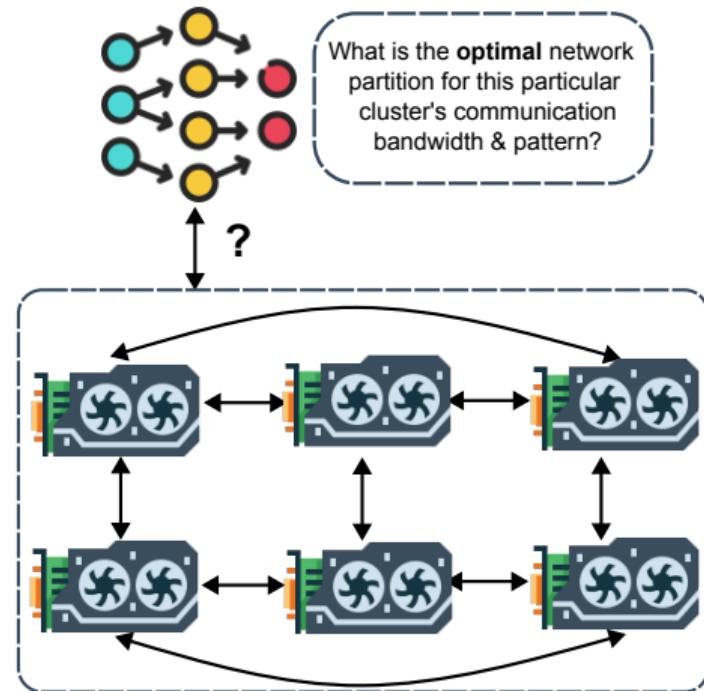
^bD. Narayanan et al., "Efficient large-scale language model training on gpu clusters using megatron-lm," in *ACM/IEEE SC*, 2021.

Figure: Envisioned adaptive system responding to network conditions

Original Vision: Adaptive Pipeline Parallelism for Multimodal Transformers

Initial Goal:

- **Adaptive pipeline parallelism system** for multimodal vision-language models (BLIP)
- **Dynamic adjustment** of pipeline configurations based on:
 - Real-time network conditions
 - Heterogeneous compute patterns
- Address inefficiencies in **static** pipeline implementations
- Extends NVIDIA's Megatron-LM framework^{a,b}
- Target: **BLIP** model on **COCO** dataset with throughput optimization



^aM. Shoeybi et al., "Megatron-Lm: Training multi-billion parameter language models using model parallelism," *arXiv preprint arXiv:1909.08053*, 2019 .

^bD. Narayanan et al., "Efficient large-scale language model training on gpu clusters using megatron-lm," in *ACM/IEEE SC*, 2021.

Figure: Envisioned adaptive system responding to network conditions

Project Pivot: From Vision to Proof-of-Concept

Challenges Encountered:

- BLIP integration with Megatron-LM proved more complex than anticipated
- Megatron-LM multimodal support still under development
- Time constraints for full implementation

Project Pivot: From Vision to Proof-of-Concept

Challenges Encountered:

- BLIP integration with Megatron-LM proved more complex than anticipated
- Megatron-LM multimodal support still under development
- Time constraints for full implementation

Strategic Pivot:

- Switched to **Llama 8B** (well-supported in Megatron-LM)
- Focused on **proof-of-concept**: Understanding P2P communication patterns
- Goal: Systematic study of TP/CP/PP/DP configurations

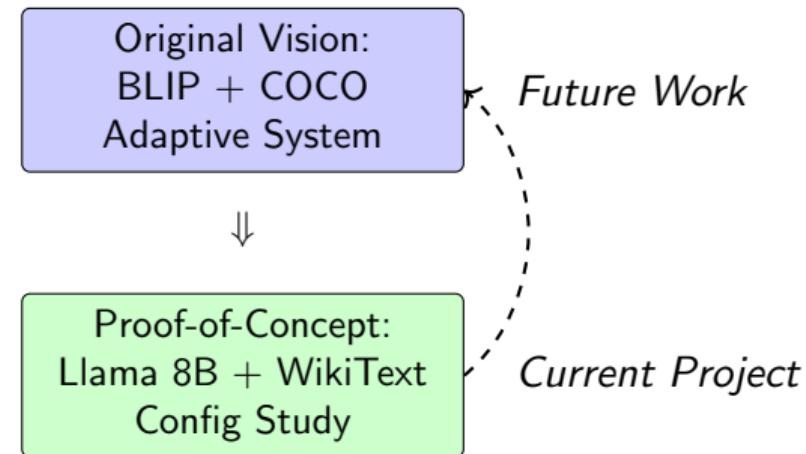


Figure: Project evolution: PoC now, full vision later

Current Project: Proof-of-Concept Study

Research Question

How do different TP/CP/PP/DP configurations affect P2P communication patterns and training throughput for large language models?

Approach

- Grid search of TP/CP/PP/DP configurations
- NCCL P2P communication profiling
- Optimize pipeline stage layout
- Basis for adaptive parallelism systems

Scope

- **Model:** LLaMA 8B (32 layers, 4096 hidden)
- **Data:** WikiText-103 (\sim 103M tokens)
- **Compute:** 4–8 GPUs, 1–2 nodes (NERSC Perlmutter)

Project Progress: What We Built

Core Infrastructure:

- Megatron-LM integration with custom P2P profiling
- LLaMA 8B model with TP/CP/PP/DP support
- WikiText-103 dataset preprocessing pipeline
- NCCL communication monitoring + profiling hooks
- Multi-node training scripts (4, 6, 8 GPUs)
- Automated logging + visualization toolkit

Project Progress: What We Built

Core Infrastructure:

- Megatron-LM integration with custom P2P profiling
- LLaMA 8B model with TP/CP/PP/DP support
- WikiText-103 dataset preprocessing pipeline
- NCCL communication monitoring + profiling hooks
- Multi-node training scripts (4, 6, 8 GPUs)
- Automated logging + visualization toolkit

Experimental Overview:

- Baseline single-GPU training runs
- TP / PP / DP / CP configuration sweeps
- P2P bandwidth profiling (intra-/inter-node)
- Training convergence + stability verification

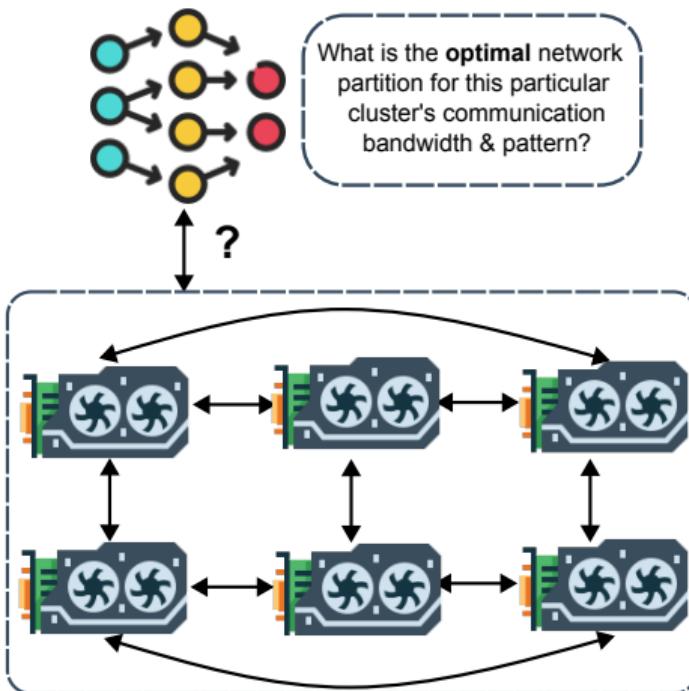
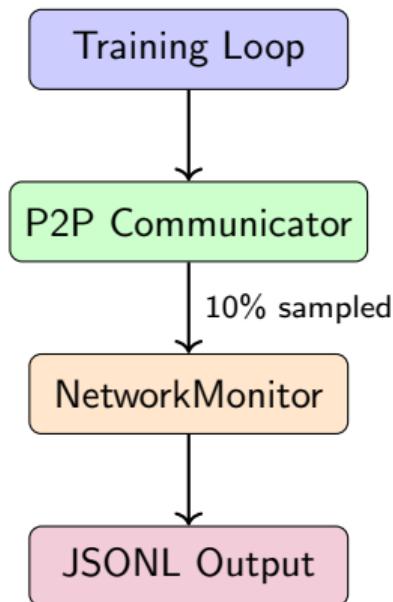


Figure: Pipeline + communication architecture

P2P Monitoring: Implementation Overview

Architecture



Output: Per-iteration JSONL with bandwidth, latency, call counts

Key Implementation Details

- **Async CUDA Events**

- Non-blocking timing via `torch.cuda.Event`
- No synchronization overhead

- **Probabilistic Sampling**

- 10% of P2P ops monitored
- Configurable sample rate

- **EMA Smoothing**

- $\alpha = 0.3$ for stable estimates
- Handles measurement noise

- **Global Registry**

- Aggregates stats across all communicators
- Thread-safe collection

P2P Monitoring: Why This Works

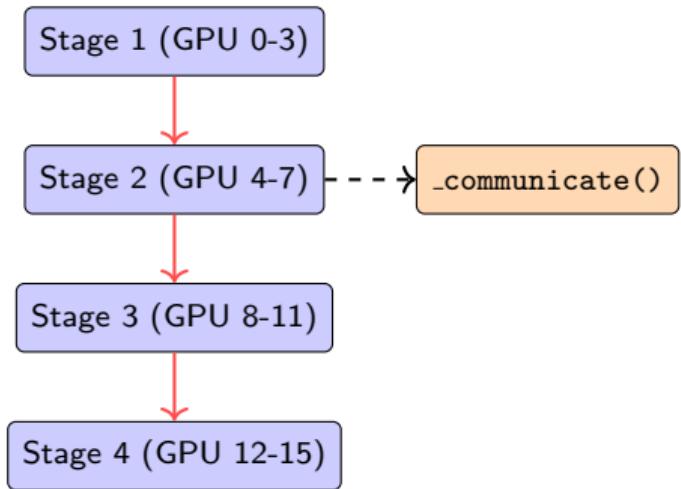
The Key Insight: Megatron-LM Has a Single Chokepoint for Pipeline Communication

How Megatron-LM Moves Data Between Pipeline Stages:

- ① Each pipeline stage needs to send activations
- ② All this communication goes through one function: `_communicate()`

Why Only Pipeline Parallelism (PP)?

- **Tensor Parallelism (TP):** Uses NCCL all-reduce *within* a node via NVLink — different code path, not P2P
- **Pipeline Parallelism (PP):** Uses point-to-point sends/recvs *between* nodes — goes through our wrapper



Red arrows =
P2P transfers
(all go through
`_communicate()`)

P2P Monitoring: How We Intercept Communication

The Wrapper Pattern: Measure Before and After Each Transfer

What We Did:

- ① Created a new class that *extends* the original communicator
- ② Start a timer → call original → stop timer
- ③ Record: how long it took, how much data moved

Pseudo-code:

```
def _communicate(self, tensor):  
    start_timer()  
    result = super().__communicate(tensor)  
    stop_timer()  
    log(time, tensor.size())  
    return result
```

Why This Approach?

- **Non-invasive:** We don't modify Megatron's core code
- **Toggleable:** A config flag switches between normal and monitored mode
- **Low overhead:** Only 10% of calls are timed (random sampling)

Files We Added:

- monitored_p2p_communication.py
- monitoring_stats_collector.py

Total: ~400 lines of code, modular design

Experimental Setup

Model Configuration

- **Architecture:** LLaMA 8B
 - 32 transformer layers
 - 32 attention heads
 - Hidden dimension: 4096
 - Vocabulary: 32,000 tokens
 - Total parameters: ~8B
- **Dataset:** WikiText-103
 - 103M tokens
 - Long-range dependencies

Note: TP = Tensor Parallel, PP = Pipeline Parallel, DP = Data Parallel, CP = Context Parallel

Experimental Setup

Model Configuration

- **Architecture:** LLaMA 8B
 - 32 transformer layers
 - 32 attention heads
 - Hidden dimension: 4096
 - Vocabulary: 32,000 tokens
 - Total parameters: ~8B
- **Dataset:** WikiText-103
 - 103M tokens
 - Long-range dependencies

9 Parallelism Configurations Tested

#	TP	CP	PP	DP	GPUs
1	4	1	4	1	16
2	2	2	4	1	16
3	1	4	4	1	16
4	1	2	8	1	16
5	1	1	16	1	16
6	2	2	2	2	16
7	4	1	2	2	16
8	2	1	8	1	16
9	1	1	8	2	16

Note: TP = Tensor Parallel, PP = Pipeline Parallel, DP = Data Parallel, CP = Context Parallel

Experimental Setup

Model Configuration

- **Architecture:** LLaMA 8B
 - 32 transformer layers
 - 32 attention heads
 - Hidden dimension: 4096
 - Vocabulary: 32,000 tokens
 - Total parameters: ~8B
- **Dataset:** WikiText-103
 - 103M tokens
 - Long-range dependencies

9 Parallelism Configurations Tested

#	TP	CP	PP	DP	GPUs
1	4	1	4	1	16
2	2	2	4	1	16
3	1	4	4	1	16
4	1	2	8	1	16
5	1	1	16	1	16
6	2	2	2	2	16
7	4	1	2	2	16
8	2	1	8	1	16
9	1	1	8	2	16

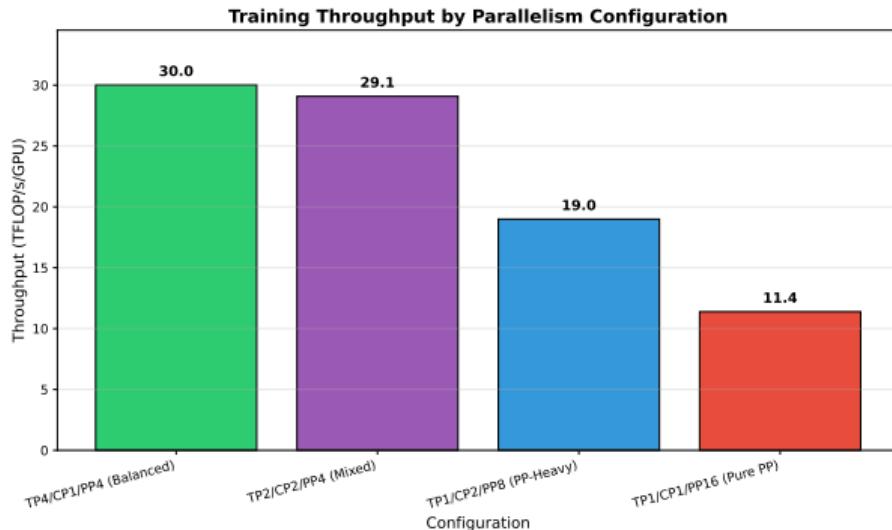
Training Hyperparameters

- Micro-batch size: 1
- Sequence length: 8192
- Learning rate: 1.5e-4
- Target iterations: 10,000

Note: TP = Tensor Parallel, PP = Pipeline Parallel, DP = Data Parallel, CP = Context Parallel

Results: Training Performance

Training Throughput by Configuration



Throughput Comparison

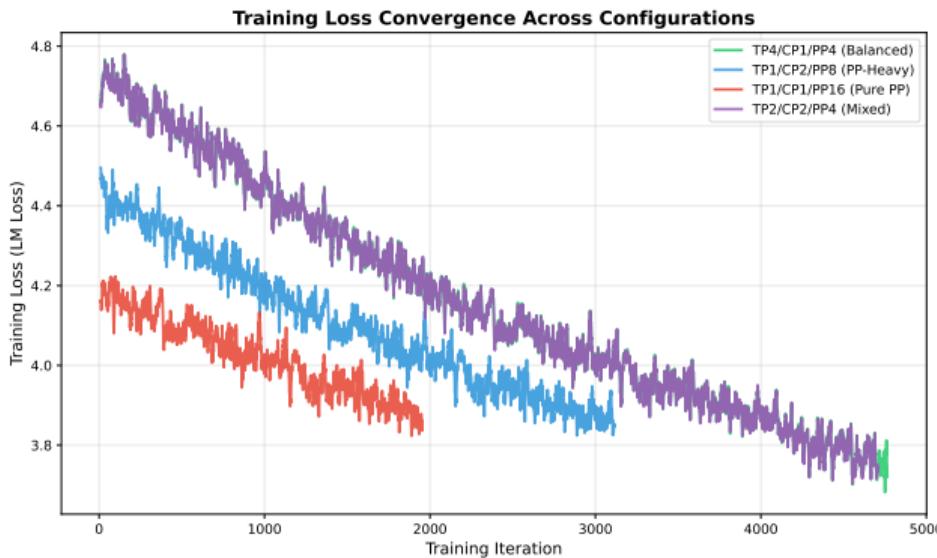
Config	TFLOP/s/GPU
TP4/CP1/PP4 (Balanced)	30.0
TP2/CP2/PP4 (Mixed)	29.1
TP1/CP2/PP8 (PP-Heavy)	19.0
TP1/CP1/PP16 (Pure PP)	11.4

Key Observations:

- **TP4/CP1/PP4** achieved highest throughput (30 TFLOP/s)
- Pure PP (PP=16) shows 62% lower throughput due to pipeline bubbles
- Balanced TP/PP configs outperform extremes

Results: Training Convergence

Training Loss Over Iterations



Convergence Validation

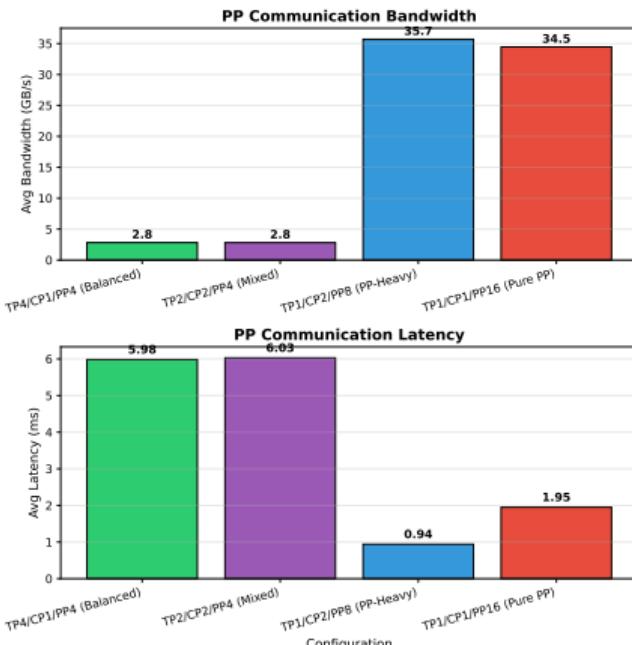
- All configurations converge to similar loss values ($\sim 3.6\text{-}4.0$)
- **Parallelism strategy does not affect model quality**—only training speed
- Loss curves show expected LLM training behavior

Final Loss Comparison

Config	Final Loss
TP4/CP1/PP4	3.65
TP1/CP1/PP16	4.00
TP1/CP2/PP8	3.92

Results: Pipeline P2P Communication Patterns

Pipeline Stage Communication Metrics



Note: Measures inter-stage PP communication only. TP uses intra-node NVLink (not monitored).

Key Findings

- **High PP (PP=8, 16):** Frequent small transfers → high BW (35 GB/s), low latency (<1ms)
- **Low PP (PP=2, 4):** Fewer, larger transfers → lower BW (2-3 GB/s), higher latency (6ms)

Trade-off Insight:

- High PP → more pipeline stages
- But: pipeline bubbles dominate → lower compute efficiency
- Optimal: Balance TP (within-node) + PP (across-node)

Results: P2P Monitoring Over Time

Bandwidth & Latency During Training



Key Findings

- **PP-heavy:** Stable 35 GB/s, ~1ms latency
- **Balanced:** Lower bandwidth (2-3 GB/s), higher compute utilization
- **Stability:** No degradation over 4000+ iterations

Implication: Predictable network behavior enables runtime-adaptive parallelism strategies

Results: Complete Configuration Comparison

9 Valid Configurations (Sorted by Throughput)

Config	TP	CP	PP	DP	TFLOP/s	Iter (s)	PP BW	PP Lat
TP4/CP1/PP4	4	1	4	1	30.0	0.89	2.8	6.0
TP2/CP2/PP4	2	2	4	1	29.1	0.91	2.8	6.0
TP1/CP4/PP4	1	4	4	1	29.0	0.93	2.8	5.9
TP1/CP1/PP8	1	1	8	2	21.2	2.43	33.5	2.0
TP1/CP2/PP8	1	2	8	1	19.0	1.36	35.7	0.9
TP2/CP1/PP8	2	1	8	1	18.9	1.37	35.3	1.0
TP4/CP1/PP2	4	1	2	2	16.7	3.07	2.8	6.1
TP2/CP2/PP2	2	2	2	2	12.1	4.24	2.7	6.2
TP1/CP1/PP16	1	1	16	1	11.4	2.24	34.5	1.9

PP BW/Lat = Pipeline parallelism inter-stage communication (GB/s, ms). TP uses NVLink (not monitored).

Key Observations:

- **Best:** TP4/CP1/PP4 balanced within-node TP with cross-node PP
- **Worst:** TP1/CP1/PP16 suffers from pipeline bubble overhead ($2.6\times$ slower)
- High PP → frequent small P2P transfers → high BW, low latency (but more bubbles)

Results: Configuration Analysis

Key Findings from 9 Experimental Runs

① Best configuration for 16 GPUs: TP=4, CP=1, PP=4

- Achieved 30.0 TFLOP/s/GPU throughput
- Balances tensor parallelism within nodes with pipeline across nodes

② Impact of pipeline depth:

- PP=4 vs PP=16: 2.6× throughput difference (30.0 vs 11.4 TFLOP/s)
- Extreme PP causes severe pipeline bubble overhead

③ Context parallelism (CP) effect:

- TP4/CP1/PP4 vs TP2/CP2/PP4: Similar throughput (~29-30 TFLOP/s)
- CP provides memory savings without major throughput penalty

④ Network topology insight:

- Intra-node: Prioritize TP (high NVLink bandwidth)
- Inter-node: Moderate PP acceptable; avoid extreme PP

General Guideline: For Perlmutter (4 GPUs/node), use **TP=4 within node, PP across nodes**

GPU Usage: Resource Allocation

Detailed GPU Hours Breakdown

Activity	Configuration	Duration	GPUs	GPU-Hours
Environment Setup	Various	4 hrs	4	16
Grid Search Runs	9 configs, 16 GPUs	1.5 hrs each	16	216
Debug & Validation	Various	8 hrs	16	128
TOTAL				~360 GPU-hours

Resource Justification

- 9 parallelism configurations tested systematically
- Each run: 1.5-hour allocation with P2P monitoring enabled
- Initial profiling adds 30-60 seconds per job for topology discovery
- All runs on NERSC Perlmutter (4 nodes × 4 A100 GPUs)

GPU Usage: Utilization Analysis

Throughput per Configuration

Config	TFLOP/s	Iter Time
TP4/CP1/PP4	30.0	890 ms
TP2/CP2/PP4	29.1	914 ms
TP1/CP4/PP4	29.0	927 ms
TP1/CP2/PP8	19.0	1365 ms
TP1/CP1/PP16	11.4	2240 ms

Performance Bottlenecks

- **PP=16:** Pipeline bubbles dominate ($2.5\times$ slower)
- **PP=8:** Moderate bubble overhead
- **TP=4, PP=4:** Best balance

P2P Monitoring Overhead

- Sample rate: 10% of P2P operations
- Measured overhead: <0.5%
- Async CUDA events: Non-blocking

PP Communication Stats

Metric	Value
Total PP P2P calls monitored	27,000+
PP BW (low PP=2,4)	2.8 GB/s
PP BW (high PP=8,16)	35 GB/s
PP latency (low PP)	6.0 ms
PP latency (high PP)	0.9 ms

*PP = pipeline stage transfers only

Lessons Learned

Key Insights from 9 Experimental Runs

- **Extreme parallelism hurts throughput** → PP=16 achieved only 11.4 TFLOP/s vs 30.0 for PP=4 ($2.6\times$ slower)
- **Pipeline bubbles dominate at high PP** → More stages = more synchronization overhead; diminishing returns beyond PP=4
- **Context Parallelism is "free"** → TP2/CP2/PP4 matched TP4/CP1/PP4 throughput while reducing memory pressure
- **Profiling overhead is negligible** → 10% sampling + async events = <0.5% overhead; enables runtime monitoring
- **Bandwidth vs efficiency trade-off** → PP-heavy configs show 35 GB/s bandwidth but poor compute utilization

Takeaway: *Balanced configs (TP=4, PP=4) outperform extremes; profiling reveals hidden bottlenecks.*

Project Completion: What We Delivered

Deliverables

- **Implementation:** Megatron-LM + LLaMA 8B, P2P profiling, multi-config scripts
- **Experiments:** Baseline + multi-GPU runs, bandwidth profiling, config comparison
- **Analysis:** Optimal config, comms bottlenecks, scaling efficiency
- **Documentation:** Open-source code + setup + reproducible training scripts

Project Completion: What We Delivered

Deliverables

- **Implementation:** Megatron-LM + LLaMA 8B, P2P profiling, multi-config scripts
- **Experiments:** Baseline + multi-GPU runs, bandwidth profiling, config comparison
- **Analysis:** Optimal config, comms bottlenecks, scaling efficiency
- **Documentation:** Open-source code + setup + reproducible training scripts

Alignment with Original Goals

Goal	Planned	PoC
Parallelism framework	BLIP	LLaMA
P2P profiling	✓	✓
Config optimization	✓	✓
Dynamic adaptation	✓	Future
Multi-node scaling	✓	✓

Project Completion: What We Delivered

Deliverables

- **Implementation:** Megatron-LM + LLaMA 8B, P2P profiling, multi-config scripts
- **Experiments:** Baseline + multi-GPU runs, bandwidth profiling, config comparison
- **Analysis:** Optimal config, comms bottlenecks, scaling efficiency
- **Documentation:** Open-source code + setup + reproducible training scripts

Alignment with Original Goals

Goal	Planned	PoC
Parallelism framework	BLIP	LLaMA
P2P profiling	✓	✓
Config optimization	✓	✓
Dynamic adaptation	✓	Future
Multi-node scaling	✓	✓

Summary

- Core goals achieved
- Model shift: BLIP → LLaMA (integration-driven)
- PoC forms base for adaptive system

Limitations & Future Directions

Limitations

- **Model scope** — LLaMA only, single size (8B)
- **Limited scale** — up to 16 GPUs / 4 nodes
- **Static configs** — fixed at training start
- **Dataset size** — WikiText-103 only

Limitations & Future Directions

Limitations

- **Model scope** — LLaMA only, single size (8B)
- **Limited scale** — up to 16 GPUs / 4 nodes
- **Static configs** — fixed at training start
- **Dataset size** — WikiText-103 only

Future Directions

- Extend to multimodal (BLIP) and multiple model scales (7B–70B)
- Scale deployment to 32–64+ GPUs across large clusters
- Implement **dynamic parallelism switching** during runtime
- Train on large corpora (COCO, RedPajama, The Pile)

Goal: *Move from a static single-model PoC to an adaptive, dynamic multimodal system.*

Future Work: From PoC to Adaptive System

Next Phase: Turn our static PoC into a data-driven, multimodal adaptive framework

- **1. Multimodal + model scaling** Integrate BLIP within our Megatron-based pipeline, characterize cross-modal latency patterns, and scale model sizes (7B–70B).
- **2. Large-cluster deployment** Port our training scripts to 32–64+ GPU environments; benchmark scaling limits and train on multi-billion-token corpora (C4, RedPajama).
- **3. Runtime adaptive control** Use our existing P2P hooks to trigger online bandwidth monitoring and dynamically switch TP/CP/PP/DP depending on network load.
- **4. Cost model + auto-configuration** Formalize trade-offs (memory, comms, bubbles) into a predictive cost model that automatically recommends optimal configurations.

Research objective: *A scalable, multimodal system that adapts its parallelism strategy in real time.*

Contributions & Impact

Research Contributions

- **P2P profiling infrastructure** NCCL hooks for Megatron-LM; reusable beyond LLaMA
- **Configuration characterization** Systematic TP/CP/PP/DP comparison + cluster-aware guidelines
- **LLaMA 8B benchmarks** Throughput + scaling efficiency as a baseline
- **Foundation for adaptive systems** PoC shows feasibility and reveals key bottlenecks

Practical Impact

Contributions & Impact

Research Contributions

- **P2P profiling infrastructure** NCCL hooks for Megatron-LM; reusable beyond LLaMA
- **Configuration characterization** Systematic TP/CP/PP/DP comparison + cluster-aware guidelines
- **LLaMA 8B benchmarks** Throughput + scaling efficiency as a baseline
- **Foundation for adaptive systems** PoC shows feasibility and reveals key bottlenecks

Practical Impact

- **For practitioners** Config selection guide, open-source implementation, profiling tools
- **For researchers** Baseline for adaptive work + communication profiling data

Open Source

- Code: GitHub Repository
- Documentation: README, examples, setup
- Training scripts: Multi-GPU ready

Toward Adaptive Multimodal Training

- **Initial objective:** Adaptive parallelism for multimodal transformers
- **Pragmatic PoC:** LLaMA 8B used to examine constraints and scalability
- **Key outcome:** Characterization of communication bottlenecks and optimal TP/CP/PP/DP configurations
- **Research value:** Infrastructure and benchmarks enabling adaptive, large-scale multimodal training

Contact: Taqiya Ehsan — te137@rutgers.edu; Khizar Anjum — ma1629@rutgers.edu